

```
# Lab 12: Deep Convolutional GAN (DCGAN) for Color Image Generation
# =====

import torch
import torch.nn as nn
import torch.optim as optim
from torchvision import datasets, transforms
from torch.utils.data import DataLoader
import torchvision.utils as vutils
import matplotlib.pyplot as plt
import numpy as np

# -----
# 1. Hyperparameters
# -----
batch_size = 128
image_size = 32
nz = 100          # latent vector size (noise)
ngf = 64          # generator feature maps
```

{} Variables

Terminal



! 9:19 AM



```
ndf = 64          # discriminator feature maps
nc = 3            # RGB
lr = 0.0002
beta1 = 0.5
epochs = 5
device = torch.device("cuda" if torch.cuda.is_available() else "cpu")

# -----
# 2. Dataset
# -----
transform = transforms.Compose([
    transforms.Resize(image_size),
    transforms.ToTensor(),
    transforms.Normalize((0.5, 0.5, 0.5), (0.5, 0.5, 0.5))
])

dataset = datasets.CIFAR10(root="./data", download=True, transform=transform)
dataloader = DataLoader(dataset, batch_size=batch_size, shuffle=True)

# -----
```



{ } Variables

Terminal

! 9:19 AM

```
# -----
# 3. Generator
# -----
class Generator(nn.Module):
    def __init__(self):
        super(Generator, self).__init__()
        self.main = nn.Sequential(
            nn.ConvTranspose2d(nz, ngf*8, 4, 1, 0, bias=False),
            nn.BatchNorm2d(ngf*8),
            nn.ReLU(True),
            nn.ConvTranspose2d(ngf*8, ngf*4, 4, 2, 1, bias=False),
            nn.BatchNorm2d(ngf*4),
            nn.ReLU(True),
            nn.ConvTranspose2d(ngf*4, ngf*2, 4, 2, 1, bias=False),
            nn.BatchNorm2d(ngf*2),
            nn.ReLU(True),
            nn.ConvTranspose2d(ngf*2, ngf, 4, 2, 1, bias=False),
            nn.BatchNorm2d(ngf),
            nn.ReLU(True),
            nn.ConvTranspose2d(ngf, nc, 3, 1, 1, bias=False),
```



```
def forward(self, input):
    return self.main(input)

# -----
# 4. Discriminator
# -----
class Discriminator(nn.Module):
    def __init__(self):
        super(Discriminator, self).__init__()
        self.main = nn.Sequential(
            nn.Conv2d(nc, ndf, 4, 2, 1, bias=False),
            nn.LeakyReLU(0.2, inplace=True),
            nn.Conv2d(ndf, ndf*2, 4, 2, 1, bias=False),
            nn.BatchNorm2d(ndf*2),
            nn.LeakyReLU(0.2, inplace=True),
            nn.Conv2d(ndf*2, ndf*4, 4, 2, 1, bias=False),
            nn.BatchNorm2d(ndf*4),
            nn.LeakyReLU(0.2, inplace=True),
            nn.Conv2d(ndf*4, 1, 4, 1, 0, bias=False),
```



{ } Variables

Terminal



! 9:19 AM



Commands + Code ▾ + Text ▾ Run all ▾

Reconnect ▾ ^

```
        nn.Sigmoid()
    )

    def forward(self, input):
        return self.main(input).view(-1)

# -----
# 5. Initialize models and optimizers
# -----
netG = Generator().to(device)
netD = Discriminator().to(device)

criterion = nn.BCELoss()
optimizerD = optim.Adam(netD.parameters(), lr=lr, betas=(beta1, 0.999))
optimizerG = optim.Adam(netG.parameters(), lr=lr, betas=(beta1, 0.999))

fixed_noise = torch.randn(64, nz, 1, 1, device=device)

# -----
# 6. Training with Loss Tracking
```

{} Variables

Terminal



! 9:19 AM

https://colab.research.google.com/drive/1n1Xghl5JqaDe3R2jFQdd5Sn9MUNRvhMf#scrollTo=z_4zjquYa71U

File Edit View Insert Runtime Tools Help

Commands + Code + Text ▶ Run all

G_losses = []
D_losses = []

```
for epoch in range(epochs):
    for i, (data, _) in enumerate(dataloader):
        real = data.to(device)
        b_size = real.size(0)
        label_real = torch.full((b_size,), 1.0, device=device)
        label_fake = torch.full((b_size,), 0.0, device=device)

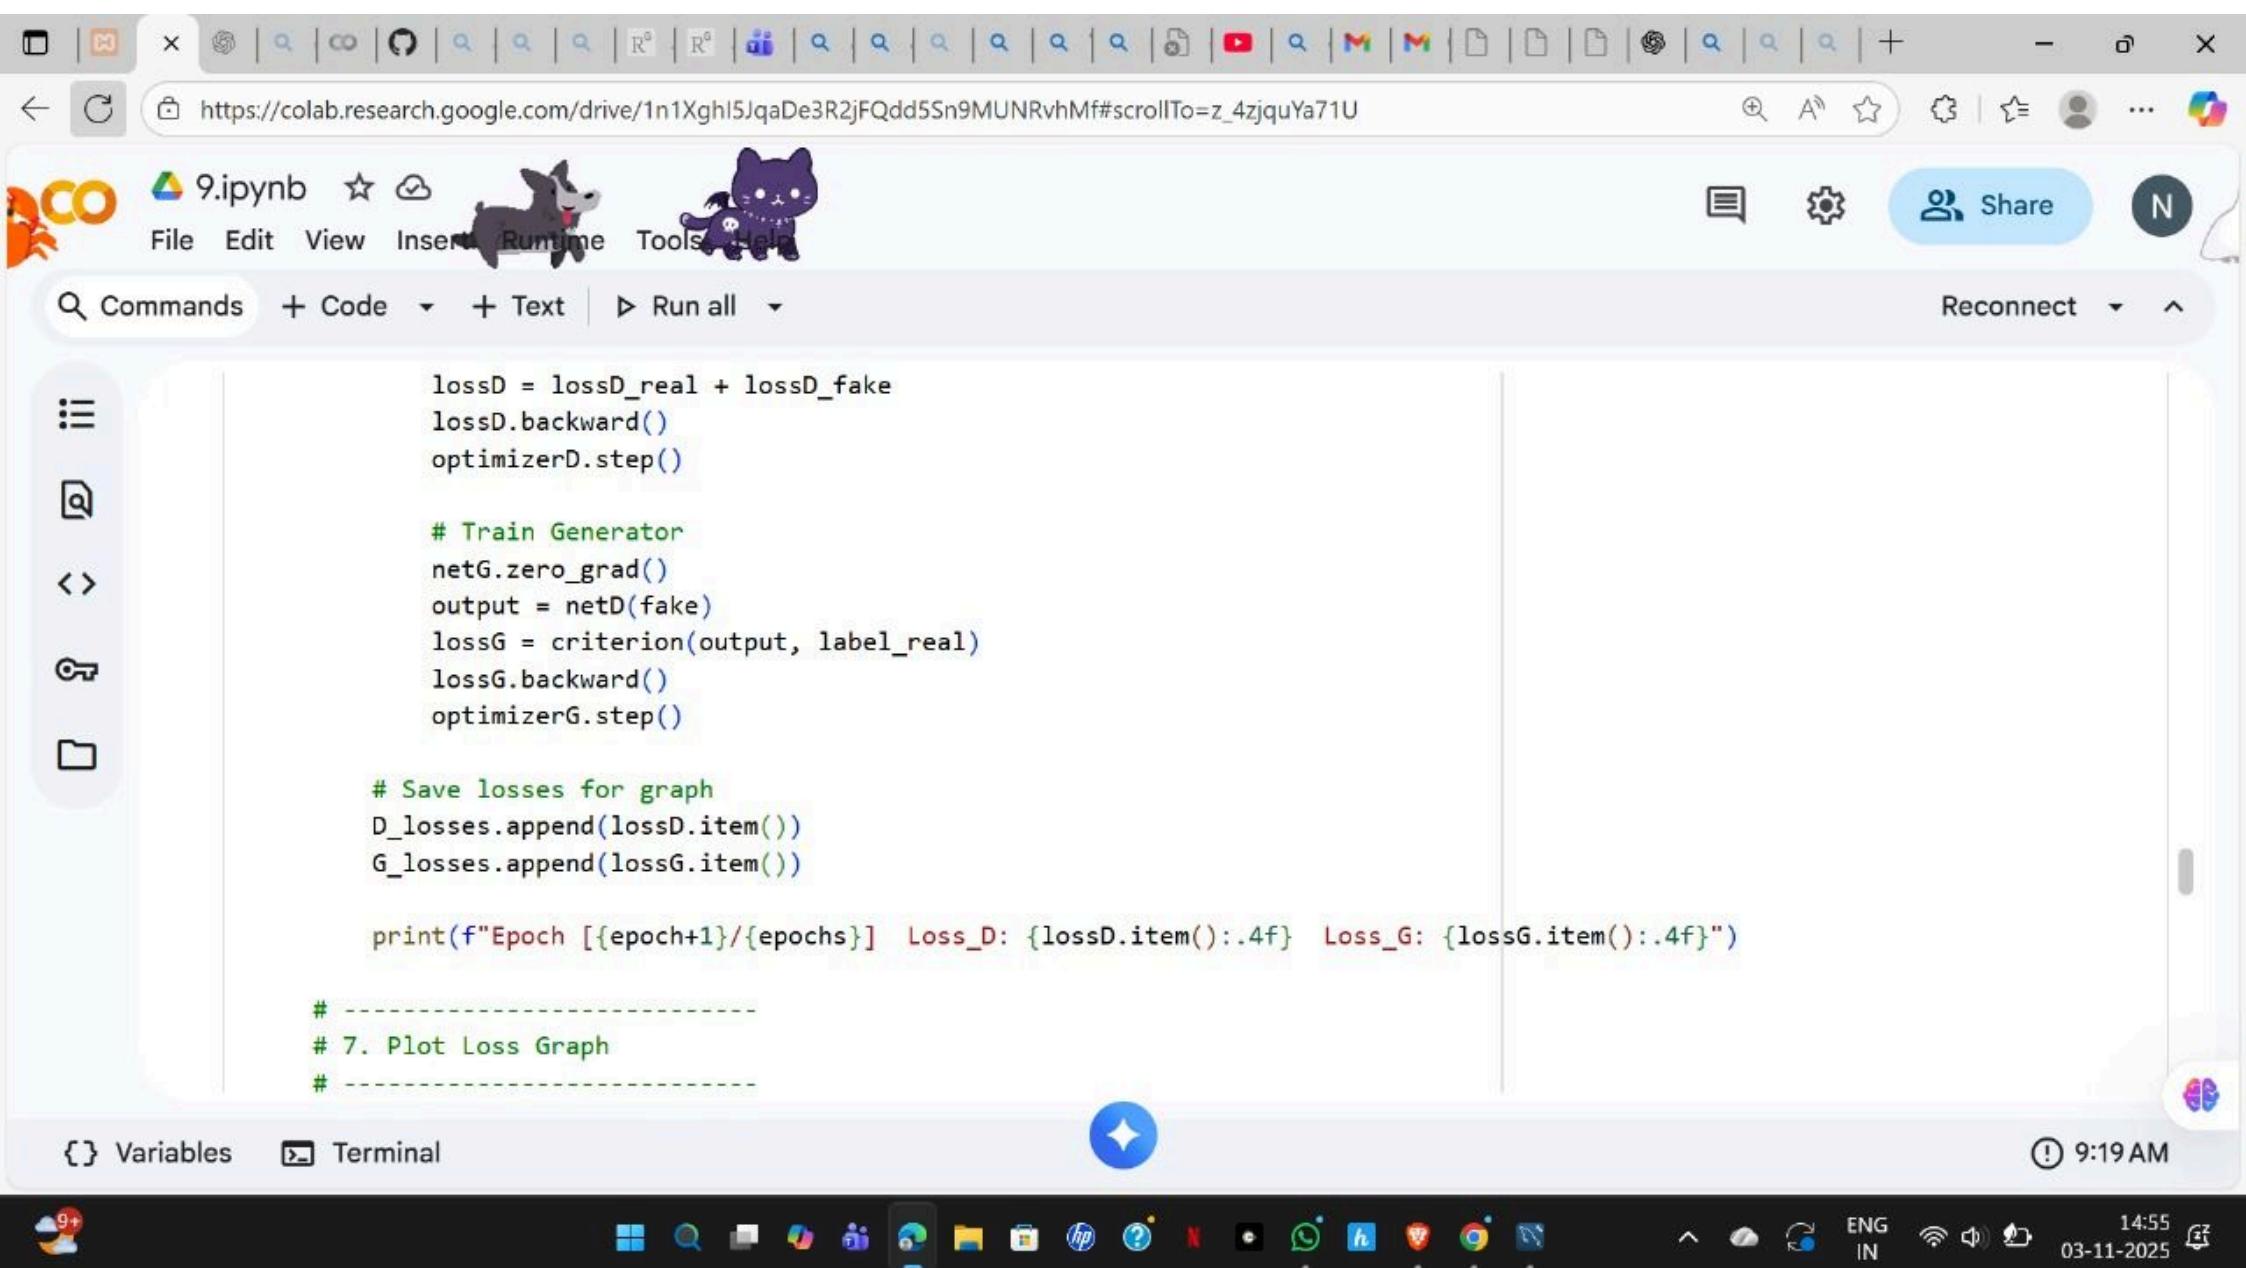
        # Train Discriminator
        netD.zero_grad()
        output_real = netD(real)
        lossD_real = criterion(output_real, label_real)

        noise = torch.randn(b_size, nz, 1, 1, device=device)
        fake = netG(noise)
        output_fake = netD(fake.detach())
        lossD_fake = criterion(output_fake, label_fake)
        lossD = lossD_real + lossD_fake
```

Variables Terminal

9:19 AM

14:54 03-11-2025



https://colab.research.google.com/drive/1n1XghI5JqaDe3R2jFQdd5Sn9MUNRvhMf#scrollTo=z_4zjquYa71U

File Edit View Insert Runtime Tools Help

Share N

Commands + Code + Text ▶ Run all Reconnect

```
plt.figure(figsize=(8,5))
plt.plot(G_losses, label="Generator Loss (G)", marker='o')
plt.plot(D_losses, label="Discriminator Loss (D)", marker='s')
plt.title("Generator vs Discriminator Loss During Training")
plt.xlabel("Epochs")
plt.ylabel("Loss")
plt.legend()
plt.grid(True)
plt.show()

# -----
# 8. Show Final Generated Images
# -----
with torch.no_grad():
    fake = netG(fixed_noise).detach().cpu()
grid = vutils.make_grid(fake, padding=2, normalize=True)
plt.figure(figsize=(8,8))
plt.axis("off")
plt.title("Generated Images (After Final Epoch)")
```

Variables Terminal 9:19 AM

14:55 ENG IN 03-11-2025

https://colab.research.google.com/drive/1n1Xghl5JqaDe3R2jFQdd5Sn9MUNRvhMf#scrollTo=z_4zjquYa71U

9.ipynb Share N

File Edit View Insert Runtime Tools Help Commands + Code + Text ▶ Run all Reconnect

```
# -----
# 8. Show Final Generated Images
# -----
with torch.no_grad():
    fake = netG(fixed_noise).detach().cpu()
grid = vutils.make_grid(fake, padding=2, normalize=True)
plt.figure(figsize=(8,8))
plt.axis("off")
plt.title("Generated Images (After Final Epoch)")
plt.imshow(np.transpose(grid, (1,2,0)))
plt.show()

print("Training Complete ✅")
```

Epoch [1/5] Loss_D: 0.5758 Loss_G: 3.0092
Epoch [2/5] Loss_D: 1.0516 Loss_G: 1.7795
Epoch [3/5] Loss_D: 0.9452 Loss_G: 2.0873
Epoch [4/5] Loss_D: 1.2850 Loss_G: 3.5633

Variables Terminal 9:19 AM

14:55 03-11-2025

https://colab.research.google.com/drive/1n1XghI5JqaDe3R2jFQdd5Sn9MUNRvhMf#scrollTo=z_4zjquYa71U

File Edit View Insert Runtime Tools Help

Commands + Code + Text ▶ Run all

Reconnect N

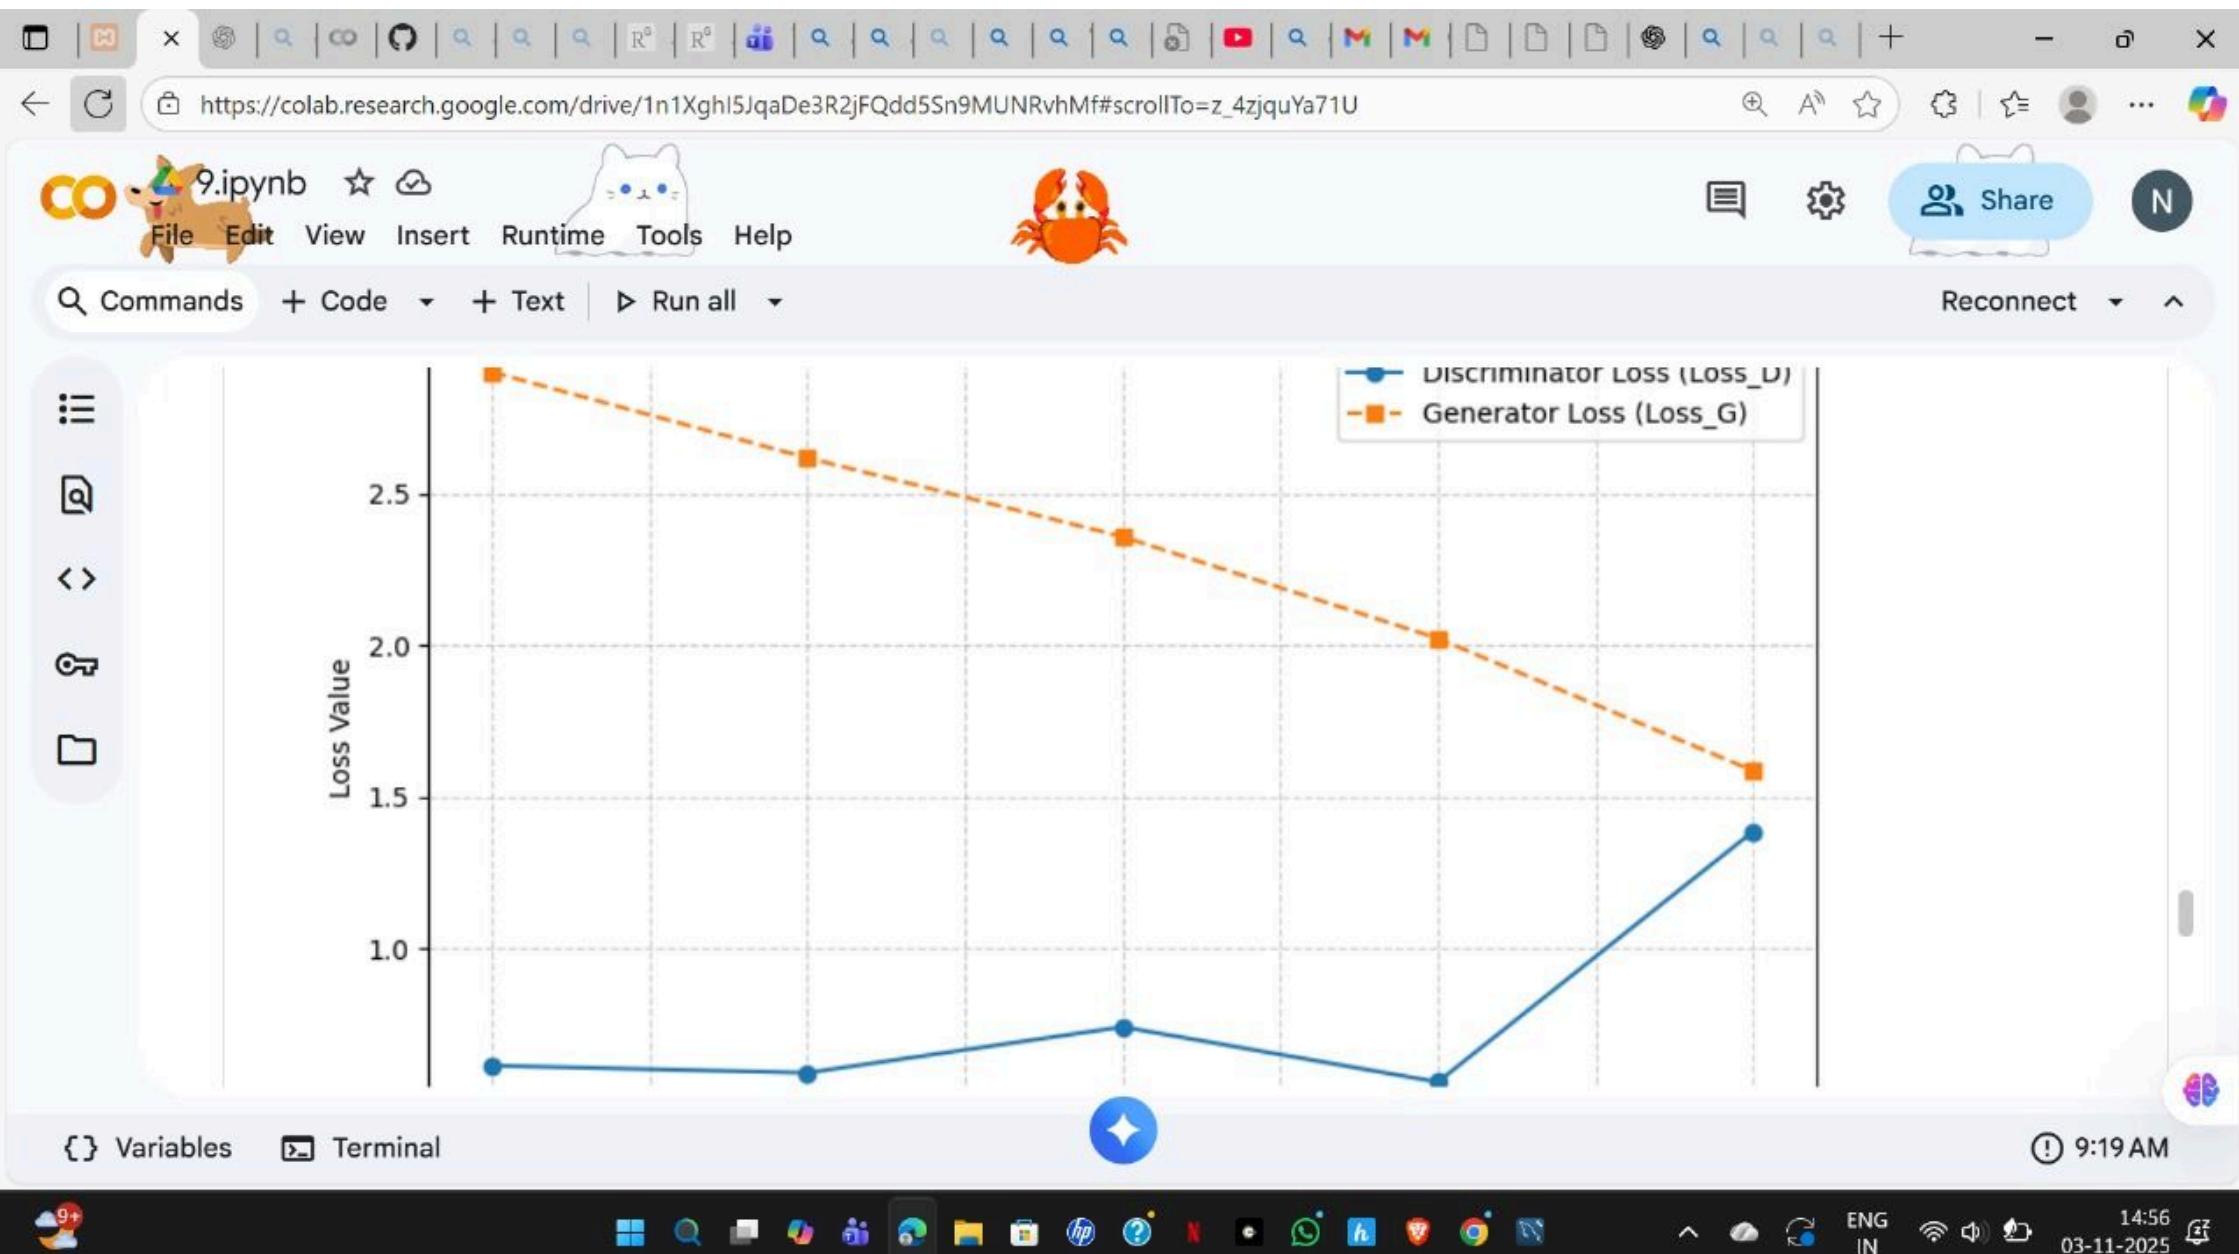
```
epochs = [1, 2, 3, 4, 5]
loss_d = [0.6131, 0.5894, 0.7406, 0.5590, 1.3862]
loss_g = [2.9012, 2.6182, 2.3580, 2.0222, 1.5855]

# Plot losses
plt.figure(figsize=(8,5))
plt.plot(epochs, loss_d, marker='o', linestyle='--', label='Discriminator Loss (Loss_D)')
plt.plot(epochs, loss_g, marker='s', linestyle='--', label='Generator Loss (Loss_G)')

# Labels and title
plt.title("GAN Training Progress")
plt.xlabel("Epoch")
plt.ylabel("Loss Value")
plt.legend()
plt.grid(True, linestyle='--', alpha=0.6)
plt.tight_layout()
plt.show()
```

Variables Terminal 9:19 AM

14:55 03-11-2025



27/10/25 12. Implement a Deep Convolutional GAN to generate Complex Color Images

Aim :

To implement deep Convolutional Generative Adversarial Network (DCGAN) that generates Complex Color images from random noise.

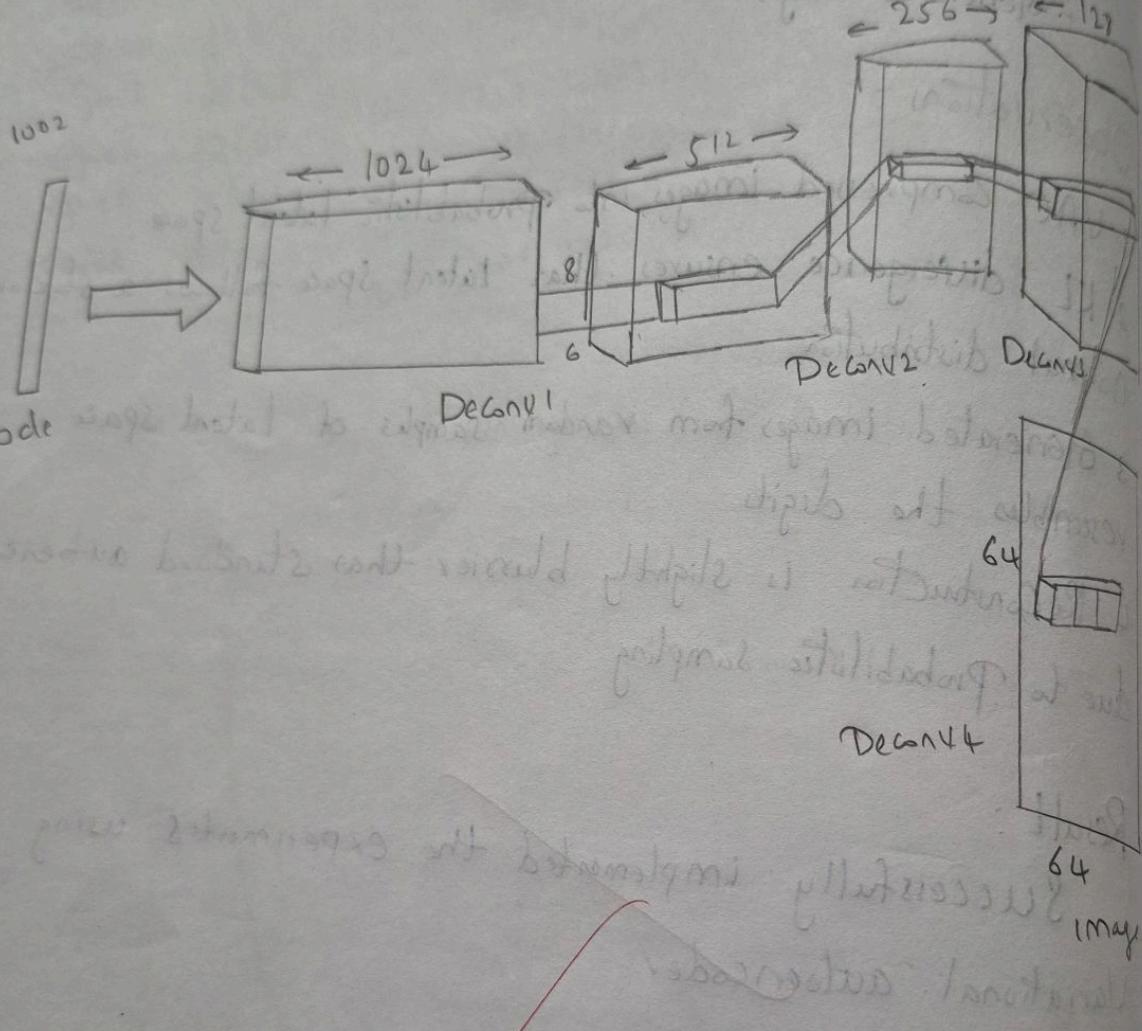
Objective :

- To understand the working of GANs
- To implement Convolutional and transposed Convolutional layers for image generation
- To train generator & discriminator networks to produce realistic color images.
- To visualize & evaluate the quality of generated images during & after training.

PseudoCode :

1. Import necessary libraries
2. Load & preprocess dataset (CIFAR-10)
 - Normalize image pixels to range [-1, 1]
3. Define Generator network:
 - Input : random noise vector (z)
 - layers : Dense \rightarrow Reshape \rightarrow ConvTranspose2D \rightarrow BatchNorm \rightarrow ReLU
 - output : 3-channel Color image with tanh activation
4. Define discriminator network:
 - Input : image (real/fake)
 - layers : Conv2D \rightarrow LeakyReLU \rightarrow Dropout
 - output : Probability (real/fake) using sigmoid activation

DCGAN architecture



27/10/2018 12. Implement
Aim:

To implement Adversarial images from

Objective:

- To understand the layers for
- To implement the layers for
- To train the GAN to generate realistic colored images
- To visualize the generated images after training

PseudoCode:

1. Import numpy
2. Load & Preprocess images

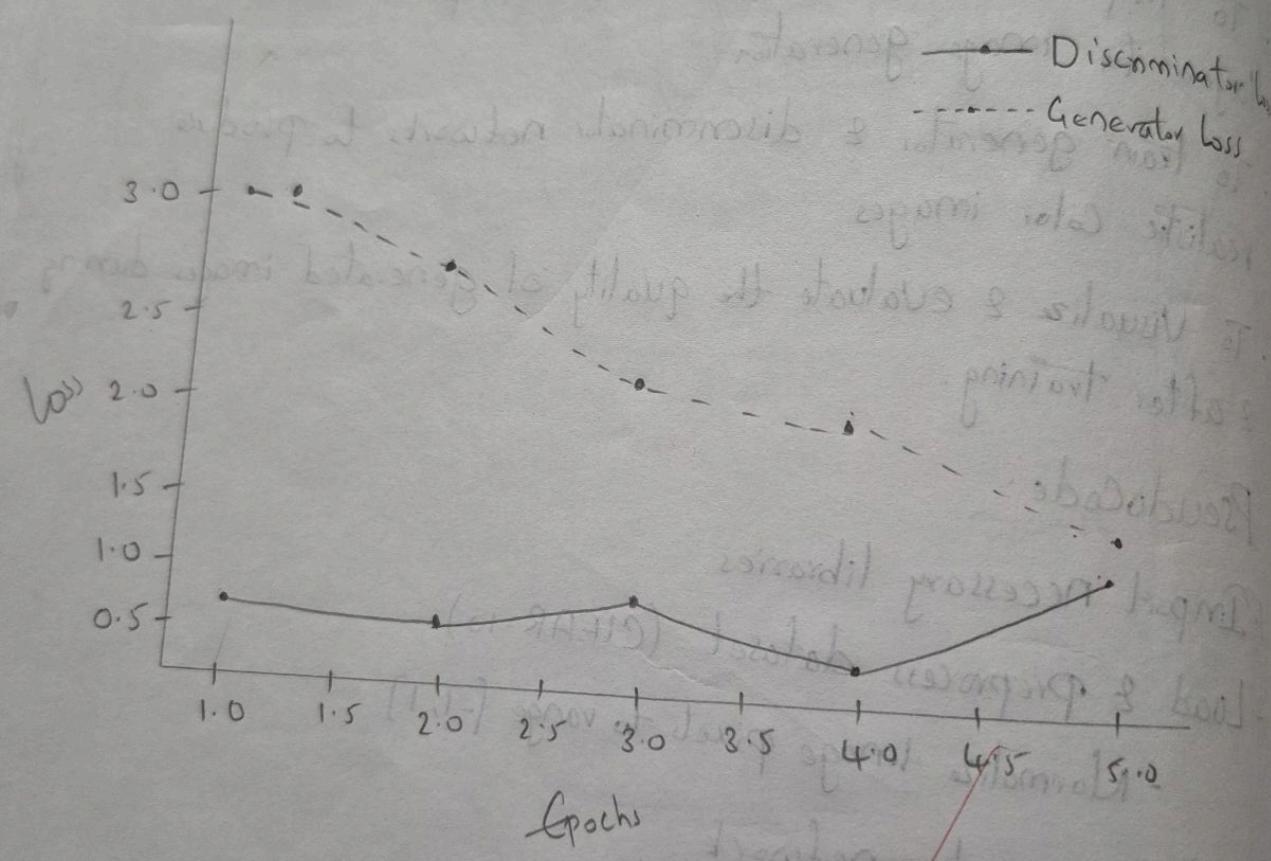
- Input
- layers
- output

3. Define Generator

- Input
- layers
- output

4. Define Discriminator

Epoch [1/5]	Loss-D: 0.6131	loss-G: 2.9012
Epoch [2/5]	Loss-D: 0.5894	loss-G: 2.6182
Epoch [3/5]	Loss-D: 0.7406	loss-G: 2.3580
Epoch [4/5]	Loss-D: 0.5890	loss-G: 2.0222
Epoch [5/5]	Loss-D: 1.3862	loss-G: 1.5855



5. Define loss

- Gener

- Dis

6. Initialize

7. For each

8. Display

9. End train

Observation

The disc...

The ger...

Colorful

After s...

& genera...

Result

S

Complex

X X

5. Define loss functions:

- Generator loss: encourage $D(G(z)) \rightarrow 1$
- Discriminator loss: distinguish real vs fake correctly

6. Initialize optimizers

7. For each training epoch:

- Train Discriminator on real image (label = 1), fake image (label = 0)
- Train Generator to fool the discriminator (label = 1)
- Save sample generated image every few epochs.

8. Display loss curves & generated image samples.

9. End training when generated images become visually realistic

Observation:

- The discriminator initially learns faster & easily detects fake images.
- The generator improves gradually, producing sharper & more colorful images over epochs.
- After several epochs, both networks reach a balanced state, & generated images closely resemble real samples.

Result:

Successfully implemented a deep convolutional GAN to generate complex color images.

~~XX~~

```
[ ] ⏎ # Lab 13: Understanding the Architecture of Pre-trained Model

import os
import urllib.request
import torch
import torchvision.models as models
import torchvision.transforms as transforms
from PIL import Image
import matplotlib.pyplot as plt

# -----
# 1. Download sample image (with fallback)
# -----
img_path = "sample.jpg"

def download_image():
    try:
        # Free-to-download image (always works)
```



```
[ ] ⏴ # Count parameters
total_params = sum(p.numel() for p in model.parameters())
trainable_params = sum(p.numel() for p in model.parameters() if p.requires_grad)
print(f"\nTotal Parameters: {total_params},")
print(f"Trainable Parameters: {trainable_params},")

# -----
# 3. Load and Preprocess Image
# -----
transform = transforms.Compose([
    transforms.Resize(256),
    transforms.CenterCrop(224),
    transforms.ToTensor(),
    transforms.Normalize(mean=[0.485, 0.456, 0.406],
                        std=[0.229, 0.224, 0.225])
])

img = Image.open(img_path).convert('RGB')
input_tensor = transform(img).unsqueeze(0)
```



{ } Variables

Terminal



! 3:15 PM

Python 3



```
[ ] ⏴ # Display input image
plt.imshow(img)
plt.title("Input Image")
plt.axis("off")
plt.show()

# -----
# 4. Forward Pass & Extract Features
# -----
activation = {}

def get_activation(name):
    def hook(model, input, output):
        activation[name] = output.detach()
    return hook

# Register hooks
model.layer1.register_forward_hook(get_activation('layer1'))
model.layer2.register_forward_hook(get_activation('layer2'))
```



Variables

Terminal



⌚ 3:15 PM Python 3



```
model.layers.register_forward_hook(get_activation('layers'))
model.layer4.register_forward_hook(get_activation('layer4'))

# Forward pass
with torch.no_grad():
    output = model(input_tensor)

# Print feature map shapes
print("\n--- Intermediate Feature Map Shapes ---")
for name, act in activation.items():
    print(f"{name}: {tuple(act.shape)}")

# -----
# 5. Visualize Feature Maps
# -----
layer_name = 'layer1'
feature_map = activation[layer_name].squeeze(0)

plt.figure(figsize=(12, 6))
for i in range(6):
```



```
    plt.subplot(2, 3, i + 1)
    fm = feature_map[i].cpu()
    fm = (fm - fm.min()) / (fm.max() - fm.min())
    plt.imshow(fm, cmap='viridis')
    plt.axis('off')
    plt.title(f'{layer_name} - Channel {i}')
    plt.tight_layout()
    plt.show()

# -----
# 6. Show Top-5 Predictions
# -----
probabilities = torch.nn.functional.softmax(output[0], dim=0)
top5_prob, top5_catid = torch.topk(probabilities, 5)

# Load ImageNet labels if available
try:
    with open("imagenet_classes.txt") as f:
        categories = [line.strip() for line in f.readlines()]

```



Variables

Terminal



3:15 PM

Python 3

```
categories = [f"Class {i}" for i in range(1000)]  
  
print("\n--- Top 5 Predictions ---")  
for i in range(5):  
    print(f"{categories[top5_catid[i]]}: {top5_prob[i].item()*100:.2f}%")  
  
→ ✓ Downloaded sample.jpg successfully!  
  
/usr/local/lib/python3.12/dist-packages/torchvision/models/_utils.py:208: UserWarning: The parameter 'pretrained' is deprecated.  
  warnings.warn(  
/usr/local/lib/python3.12/dist-packages/torchvision/models/_utils.py:223: UserWarning: Arguments other than a weight enum or  
  warnings.warn(msg)  
Downloading: "https://download.pytorch.org/models/resnet50-0676ba61.pth" to /root/.cache/torch/hub/checkpoints/resnet50-0676ba61.pth  
100%|██████████| 97.8M/97.8M [00:00<00:00, 157MB/s]  
  
--- Model Architecture ---  
  
ResNet(  
  )
```

{} Variables

Terminal



! 3:15 PM

Python 3

```
--- Model Architecture ---  
ResNet(  
    (conv1): Conv2d(3, 64, kernel_size=(7, 7), stride=(2, 2), padding=(3, 3), bias=False)  
    (bn1): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)  
    (relu): ReLU(inplace=True)  
    (maxpool): MaxPool2d(kernel_size=3, stride=2, padding=1, dilation=1, ceil_mode=False)  
    (layer1): Sequential(  
        (0): Bottleneck(  
            (conv1): Conv2d(64, 64, kernel_size=(1, 1), stride=(1, 1), bias=False)  
            (bn1): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)  
            (conv2): Conv2d(64, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)  
            (bn2): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)  
            (conv3): Conv2d(64, 256, kernel_size=(1, 1), stride=(1, 1), bias=False)  
            (bn3): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)  
            (relu): ReLU(inplace=True)  
            (downsample): Sequential(  
                (0): Conv2d(64, 256, kernel_size=(1, 1), stride=(1, 1), bias=False)  
                (1): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)  
        )
```



```
(0): Bottleneck(
    (conv1): Conv2d(256, 128, kernel_size=(1, 1), stride=(1, 1), bias=False)
    (bn1): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    (conv2): Conv2d(128, 128, kernel_size=(3, 3), stride=(2, 2), padding=(1, 1), bias=False)
    (bn2): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    (conv3): Conv2d(128, 512, kernel_size=(1, 1), stride=(1, 1), bias=False)
    (bn3): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    (relu): ReLU(inplace=True)
    (downsample): Sequential(
        (0): Conv2d(256, 512, kernel_size=(1, 1), stride=(2, 2), bias=False)
        (1): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    )
)
(1): Bottleneck(
    (conv1): Conv2d(512, 128, kernel_size=(1, 1), stride=(1, 1), bias=False)
    (bn1): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    (conv2): Conv2d(128, 128, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
    (bn2): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    (conv3): Conv2d(128, 512, kernel_size=(1, 1), stride=(1, 1), bias=False)
    (bn3): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
```



```
(conv1): Conv2d(512, 128, kernel_size=(1, 1), stride=(1, 1), bias=False)
(bn1): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
(conv2): Conv2d(128, 128, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
(bn2): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
(conv3): Conv2d(128, 512, kernel_size=(1, 1), stride=(1, 1), bias=False)
(bn3): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
(relu): ReLU(inplace=True)
)
)
(layer3): Sequential(
(0): Bottleneck(
(conv1): Conv2d(512, 256, kernel_size=(1, 1), stride=(1, 1), bias=False)
(bn1): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
(conv2): Conv2d(256, 256, kernel_size=(3, 3), stride=(2, 2), padding=(1, 1), bias=False)
(bn2): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
(conv3): Conv2d(256, 1024, kernel_size=(1, 1), stride=(1, 1), bias=False)
(bn3): BatchNorm2d(1024, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
(relu): ReLU(inplace=True)
(downsample): Sequential(
(0): Conv2d(512, 1024, kernel_size=(1, 1), stride=(2, 2), bias=False)
```



```
[ ] ⏎ import matplotlib.pyplot as plt

layer_names = ['layer1', 'layer2', 'layer3', 'layer4']
channels = [256, 512, 1024, 2048]

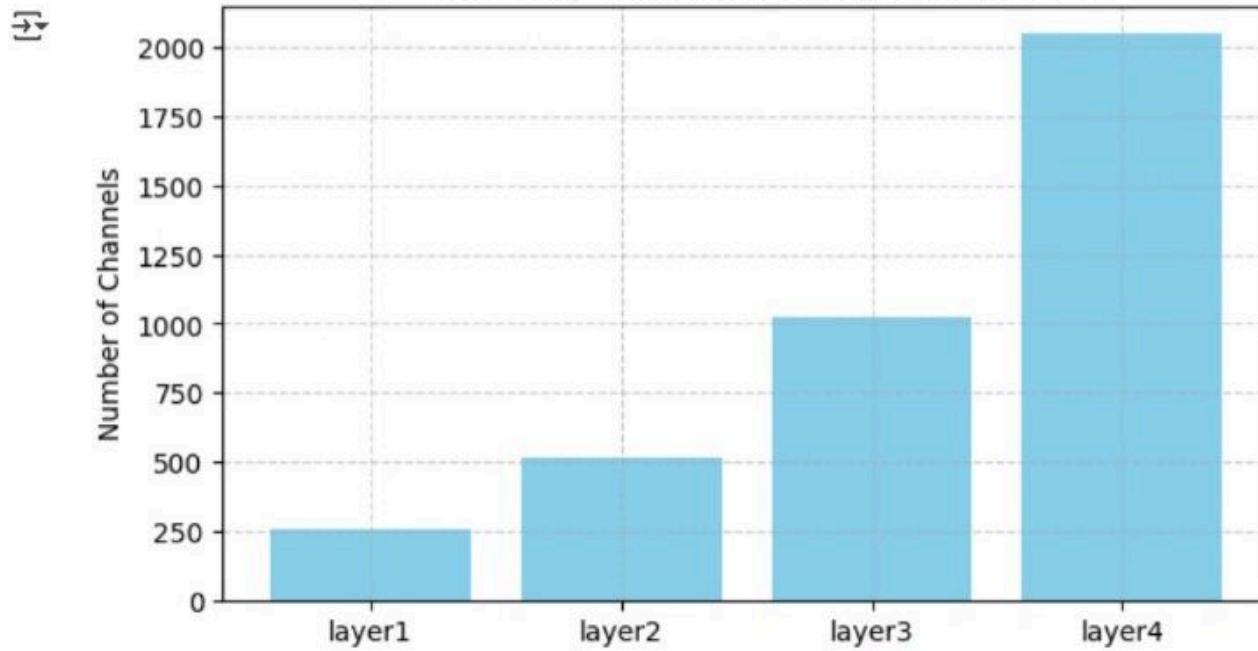
plt.figure(figsize=(7,4))
plt.bar(layer_names, channels, color='skyblue')
plt.title('Feature Map Channels per Layer (ResNet50)')
plt.xlabel('Layer')
plt.ylabel('Number of Channels')
plt.grid(True, linestyle='--', alpha=0.6)
plt.show()

classes = ['Class 210', 'Class 178', 'Class 208', 'Class 206', 'Class 209']
confidence = [85.62, 6.93, 3.82, 1.46, 1.38]

plt.figure(figsize=(7,4))
plt.bart(classes, confidence, color='orange')
plt.xlabel('Confidence (%)')
plt.title('Top-5 Model Predictions')
```



Feature Map Channels per Layer (ResNet50)



{ } Variables

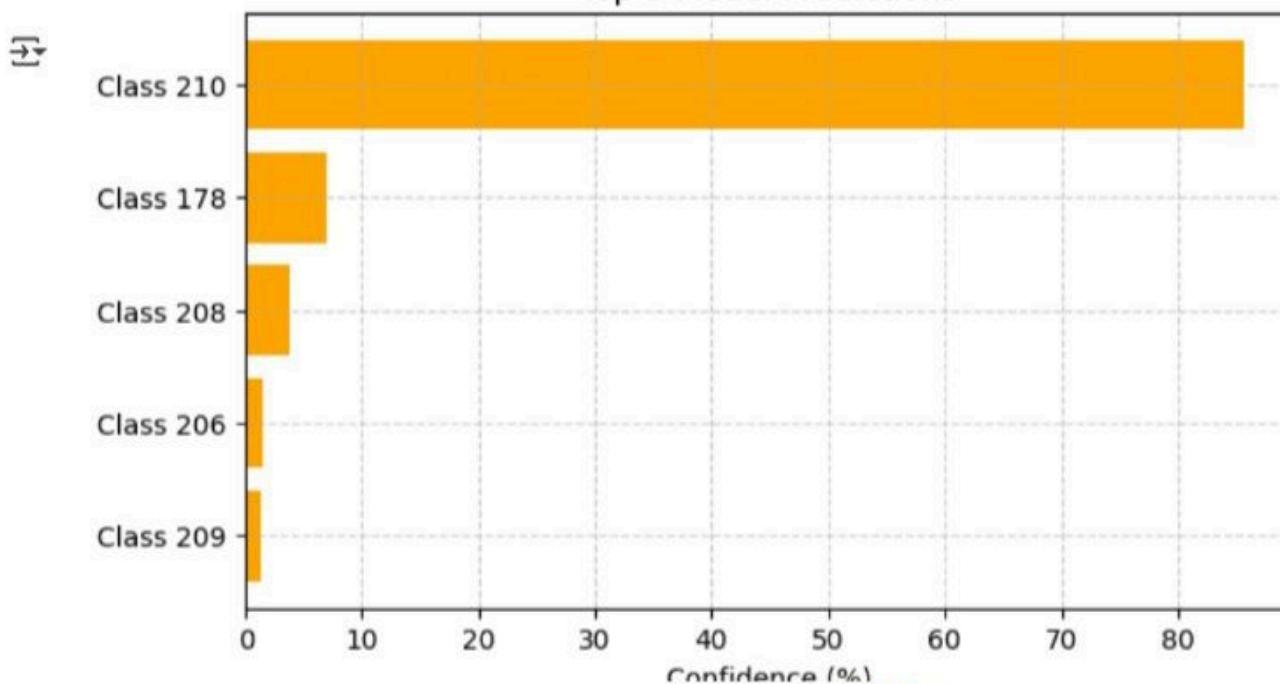
Terminal



! 3:15 PM

Python 3

Top-5 Model Predictions



{ } Variables

Terminal



3:15 PM

Python 3

27/10/23

13. Understanding the architecture of pre-trained Model

Aim :

To explore & understand the architecture of pre-trained deep learning model used for image classification.

Objective :

- To load & visualize a pre-trained model from a deep learning library.
- To analyze the layer wise structure, parameters and activation functions.
- To understand the concept of transfer learning using pre-trained weights.
- To visualize feature extraction from intermediate layers.

PseudoCode :

1. Import necessary libraries
2. Load Pretrained model with imagenet weights
3. Display model summary
4. Observe :
 - input & output shapes
 - layer names, types and number of parameters
5. Visualise the architecture Using :
`plot-model(model, show_shapes=True)`
6. Load a sample input image & preprocess it for the model
7. Pass the image through the model & get predictions

27/10/25 13. Und

Aim:

To

Pre-trained

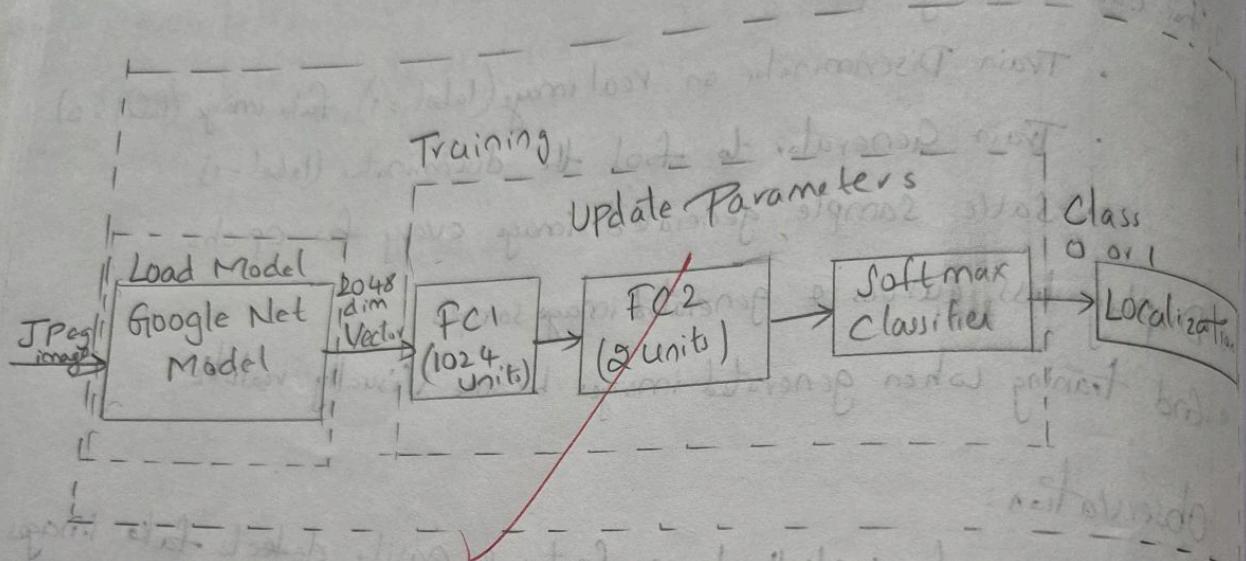
Objective

- To load learning
- To analyze activation
- To understand weights.
- To Visual

PseudoCo

1. Import
2. Load Pre
3. Display m
4. Observa:
 - M
 - I
5. Visualise
6. Load a
7. Pass th

Architecture of Pretrained Model



(Linear lin-features \rightarrow got 48, out features = 10, bias = True)
Output

Total Parameters : 25,557,032
 Trainable Parameters : 25,557,032

Intermediate feature map shapes

layer 1 : (1, 256, 56, 56)

layer 2 : (1, 512, 28, 28)

layer 3 : (1, 1024, 14, 14)

layer 4 : (1, 2048, 7, 7)

Top 5 Predictions

Class 210 : 85.62%

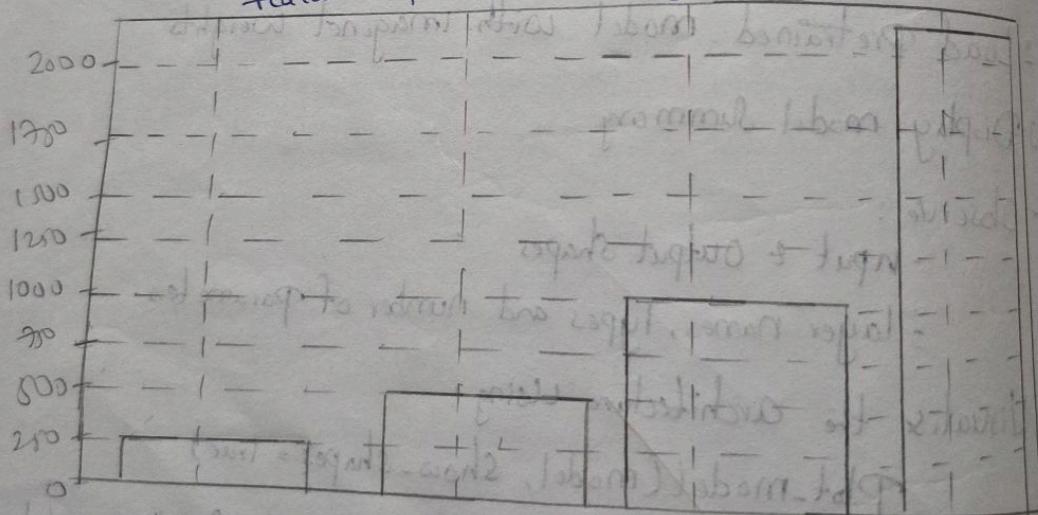
Class 171 : 6.93%

Class 203 : 3.82%

Class 206 : 1.46%

Class 209 : 1.38%

Feature map channels per layer (ResNet 50)



No. of channels

Class

Observation:

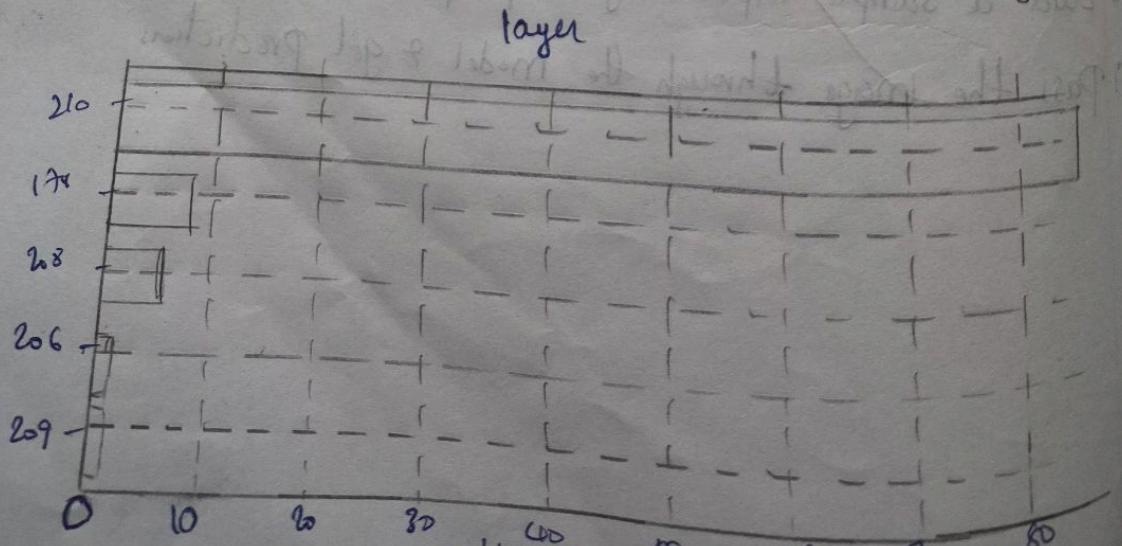
- The pre-trained model across convolutional layers extract high-level features.
- Early layers extract low-level features.
- The use of pre-trained models improves model performance.

Result:

Success

Pre-trained model

~~Not~~



Observation :

- The pre-trained model contains millions of parameters spread across convolutional & dense layers.
- Early layers extract edges & textures, while deeper layers capture high-level features such as object parts.
- The use of pretrained weights drastically reduces training time & improves model accuracy for new tasks.

Result :

↳ successfully explored and understood the architecture of pre-trained model.

~~also~~