

14. Implement a pre-trained CNN Model as Feature Extractor.

Aim:

To implement a pre-trained CNN as feature extractor using transfer learning



Objective:

- To understand how a CNN pre-trained on ImageNet can extract visual features
- To freeze Convolutional layers and use their output features for new tasks.
- To train only the classifier head for faster and more accurate learning
- To observe how transfer learning improves performance with limited data.

Pseudo Code:

1. Import `torchvision.models` and `torch` libraries
2. Load a Pretrained CNN from `torchvision`
3. Freeze all Convolutional layers to prevent retraining
4. Replace the final classifier layer for new number of classes
5. Load a small dataset (eg, CIFAR-10)
6. Extract features using the pretrained backbone
7. Train only the new classifier head
8. Evaluate accuracy and visualize results.

Pre-trained CNN model Architecture



**Freeze Pretrained CNN architecture
layers**

Retain

Observation

| Aspect | ResNet18 Feature Extractor | Training from scratch |
|---------------|------------------------------------|---------------------------|
| Training time | Very low | High |
| Accuracy | ~85%. | ~65%. |
| Required data | small | large |
| Key benefit | Reuses pre-trained visual features | Must learn everything new |


Result :

Successfully implemented a pre-trained CNN as a feature extractor

~~Pre-trained~~

Output

Epoch [1/2], loss: 0.8107
 Epoch [2/2], loss: 0.6423



Test accuracy: 79.25%.

[1]

```
# Lab 14: Implement a Pre-trained CNN model as a Feature Extractor
# =====

import torch
import torch.nn as nn
import torch.optim as optim
from torchvision import models, datasets, transforms
from torch.utils.data import DataLoader
import matplotlib.pyplot as plt
import os

# -----
# Step 1: Define Transformations
# -----
transform = transforms.Compose([
    transforms.Resize((224, 224)),
    transforms.ToTensor(),
])
```



[1]

```
# Download a small example dataset (CIFAR-10)
train_data = datasets.CIFAR10(root='./data', train=True, download=True, transform=transform)
test_data = datasets.CIFAR10(root='./data', train=False, download=True, transform=transform)

train_loader = DataLoader(train_data, batch_size=32, shuffle=True)
test_loader = DataLoader(test_data, batch_size=32, shuffle=False)

# -----
# Step 2: Load Pre-trained Model
# -----
model = models.resnet18(pretrained=True)

# Freeze all convolutional layers
for param in model.parameters():
    param.requires_grad = False

# -----
# Step 3: Replace the Classifier
```



Commands + Code ▾ + Text ▾ Run all ▾ ✓ RAM Disk

[1] num_features = model.fc.in_features
model.fc = nn.Linear(num_features, 10) # CIFAR-10 has 10 classes

Move model to device
device = torch.device("cuda" if torch.cuda.is_available() else "cpu")
model = model.to(device)

Step 4: Define Loss and Optimizer

criterion = nn.CrossEntropyLoss()
optimizer = optim.Adam(model.fc.parameters(), lr=0.001)

Step 5: Train Only the Classifier

num_epochs = 2 # keep small for quick run
train_loss_list = []

Commands + Code ▾ + Text ▾ Run all ▾

✓ RAM Disk

```
[1]  for epoch in range(num_epochs):
        running_loss = 0.0
        for images, labels in train_loader:
            images, labels = images.to(device), labels.to(device)

            optimizer.zero_grad()
            outputs = model(images)
            loss = criterion(outputs, labels)
            loss.backward()
            optimizer.step()
            running_loss += loss.item()

        avg_loss = running_loss / len(train_loader)
        train_loss_list.append(avg_loss)
        print(f"Epoch [{epoch+1}/{num_epochs}], Loss: {avg_loss:.4f}")

print("\n✓ Training completed successfully!")
```



```
[1] plt.plot(train_loss_list, marker='o')
plt.title("Training Loss per Epoch")
plt.xlabel("Epoch")
plt.ylabel("Loss")
plt.grid(True)
plt.show()

# -----
# Step 7: Test Accuracy
# -----

correct = 0
total = 0
model.eval()
with torch.no_grad():
    for images, labels in test_loader:
        images, labels = images.to(device), labels.to(device)
        outputs = model(images)
        _, predicted = torch.max(outputs, 1)
        total += labels.size(0)
```



[1]

```
    _, predicted = torch.max(outputs, 1)
    total += labels.size(0)
    correct += (predicted == labels).sum().item()

    accuracy = 100 * correct / total
    print(f"\n Test Accuracy: {accuracy:.2f}%")
```



```
→ 100%|██████████| 170M/170M [00:10<00:00, 16.0MB/s]
/usr/local/lib/python3.12/dist-packages/torchvision/models/_utils.py:208: UserWarning: The parameter 'pretrained' is deprecated
  warnings.warn(
/usr/local/lib/python3.12/dist-packages/torchvision/models/_utils.py:223: UserWarning: Arguments other than a weight enum or
  warnings.warn(msg)
Downloading: "https://download.pytorch.org/models/resnet18-f37072fd.pth" to /root/.cache/torch/hub/checkpoints/resnet18-f3707
100%|██████████| 44.7M/44.7M [00:00<00:00, 156MB/s]
Epoch [1/2], Loss: 0.8107
Epoch [2/2], Loss: 0.6423
```

Training completed successfully!



↑ ↓ ⚪ 🗑️ ⏮


☰ 🔍 < > 🔑 📁

{ } Variables ☐ Terminal



⌚ 3:15 PM Python 3





Test Accuracy: 79.21%



15. Implement a YOLO Model to detect object

Aim :

To implement YOLO Model for real time object detection and understand its working Principle

Objective :

- To study how YOLO detects multiple objects in single image
- To understand the architecture and working of a pre-trained YOLO model.
- To use transfer learning for custom object detection
- To visualize the bounding boxes and class labels Predicted by YOLO.

Pseudo Code :

1. Install and import ultralytics
2. Load a Pre-trained YOLO model
3. Load a test image or use a camera frame
4. Run the Model.predict() method to detect Objects
5. Display bounding boxes and class labels
6. Save the annotated output image

Observation :

1. The Pre-trained YOLOv8s model successfully detected multiple Objects such as cars, buses & persons in a single frame.
2. Each object was enclosed in a bounding box with a class label and confidence score.

3. The inference time per image was very low, showing YOLO's high efficiency for real-time use
4. The model demonstrated strong generalization without additional training.
5. The visualization clearly showed YOLO's ability to detect overlapping objects in complex scenes

Result :

~~successfully implemented a pre-trained YOLO model.~~

~~STL~~

Output

Image 111: 720 x 1280 & person, 1 tree, 1 cell phone
Speed 89.3ms Preprocess, 822.9ms inference, 45.7ms NMs
Per image at shape (1, 3, 384, 640)

- 3. The high
- 4. The tra
- 5. The Ob
- R

```
print("\n Loading YOLOv5 model (this may take a few seconds)...")
model = torch.hub.load('ultralytics/yolov5', 'yolov5s', pretrained=True)

# Step 3: Perform inference on the image
img = Image.open(img_path)
results = model(img)

# Step 4: Print detection results
print("\n Detection Results")
results.print() # prints detected objects, confidence, coordinates

# Step 5: Display the image with bounding boxes
results.show() # opens the image window with labels and boxes

# Step 6: (Optional) Save output image to a folder
output_folder = "yolo_output"
results.save(save_dir=output_folder)
print(f"\n Detection completed! Output saved in folder: '{output_folder}'")

# -----
```



{ } Variables

Terminal

! 9:19 AM

```
detected_img_path = os.path.join(output_folder, os.path.basename(img_path))
if os.path.exists(detected_img_path):
    detected_img = Image.open(detected_img_path)
    plt.imshow(detected_img)
    plt.axis('off')
    plt.title("YOLOv5 Object Detection Result")
    plt.show()
```




→ Image already exists: street.jpg

```
Loading YOLOv5 model (this may take a few seconds)...
Creating new Ultralytics Settings v0.0.6 file ✓
View Ultralytics Settings with 'yolo settings' or at '/root/.config/Ultralytics/settings.json'
Update Settings with 'yolo settings key=value', i.e. 'yolo settings runs_dir=path/to/dir'. For help see https://docs.ultralytics.com
Using cache found in /root/.cache/torch/hub/ultralytics_yolov5_master
YOLOv5 🚀 2025-11-1 Python-3.12.12 torch-2.8.0+cu126 CPU
```



```
Image already exists: street.jpg
Loading YOLOv5 model (this may take a few seconds)...
creating new ultralytics settings v5.0.6 file 📁
View Ultralytics Settings with 'yolo settings' or at '/root/.config/ultralytics/settings.json'
Update Settings with 'yolo settings key=value', i.e. 'yolo settings runs_dir=path/to/dir'. For help see https://docs.ultralytics.com/quickstart/ultralytics-settimes.
using cache found in /root/.cache/torch/hub/ultralytics_yolov5s_master
YOLOv5s # 2025-11-1 PyTorch-3.12.12 torch-2.8.0+cu126 CPU
downloading https://ultralytics.com/releases/yolov5s.pt to yolov5s.pt...
size: [██████████] 14.1M/14.1M [00:00<00:00, 114MB/s]

Fusing layers...
YOLOv5s summary: 213 layers, 7225885 parameters, 0 gradients, 16.4 GFLOPs
Adding AutoShape...
/root/.cache/torch/hub/ultralytics_yolov5s_master/models/common.py:986: FutureWarning: 'torch.cuda.amp.autocast(args...)' is deprecated. Please use 'torch.amp.autocast('cuda', args...)' instead.
with amp.autocast(autocast):
Image 1/1: 🕵️ Q 🎯 S - P 🚶 1 phone
Speed: 89... 🕵️ Q 🎯 S - P 🚶 1 phone, 46.7ms RMS per image at shape (1, 3, 384, 640)

detection results

```